

# Math 300 NTI Lesson 6

group\_by, mutate, and arrange

Professor Bradley Warner

June, 2022

## Contents

Objectives . . . . .	1
Reading . . . . .	1
Lesson . . . . .	1
Documenting software . . . . .	8

## Objectives

1. Use the `group_by()` function to create aggregated data frames to use with other functions, in particular `summarize()`, to explore, explain, and visualize.
2. Use the `mutate()` function to create new variables in a data frame in order to explore, explain, and visualize.
3. Use the `arrange()` function to sort data frames to explore, explain, and visualize.

## Reading

Chapter 3.4 - 3.6

## Lesson

Remember that you will be running this more like a lab than a lecture. You want them using R and answering questions. Have them open the notes rmd and work through it together.

Work through the learning checks LC3.5 - LC3.12.

- We changed the scaffolded code method. We have `eval=FALSE` so that R does not try to evaluate the code chunk. They have to remove this and then complete the code.
- It is important to note that the `group_by()` function doesn't change data frames by itself. Rather it changes the meta-data, or data about the data, specifically the grouping structure. It is only after we apply the `summarize()` function that the data frame changes. The book does a good job explaining meta-data.

- The `group_by()` can be used on more than two variables but they must be in the same call to `group_by()`.
- Using `arrange()` is straight forward expect the use of `desc()` within the `arrange()` call to sort in decreasing order.
- As a rough rule of thumb, as long as you are not losing original information that you might need later, it's acceptable practice to overwrite existing data frames with updated ones.
- LC 3.6 is difficult. Let them explore and wrestle with this question. The warning can be ignored. We will not experiment with the `.groups` option.
- LC 3.12 is more difficult as we combined code. The default use of `geom_boxplot()` works for exploring data but we provided code on how to clean up the x-axis. Discuss this code if you want and have time.
- The use of `kable()` is only to have the output printed in a form that looks good. This is not something we need to present to the students.

## Setup

```
library(nycflights13)
library(ggplot2)
library(dplyr)
```

## LC 3.5 (Objective 1)

(LC3.5) Recall from Chapter 2 when we looked at plots of temperatures by months in NYC. What does the standard deviation column in the `summary_monthly_temp` data frame, which we need to create from the code at the section 3.4, tell us about temperatures in New York City throughout the year?

**Solution:**

```
# Code from the book
summary_temp_by_month <- weather %>%
  group_by(month) %>%
  summarize(
    mean = mean(temp, na.rm = TRUE),
    std_dev = sd(temp, na.rm = TRUE)
  )
```

```
# Output
summary_temp_by_month
```

```
## # A tibble: 12 x 3
##   month mean std_dev
##   <int> <dbl> <dbl>
## 1     1  35.6  10.2
## 2     2  34.3   6.98
## 3     3  39.9   6.25
## 4     4  51.7   8.79
## 5     5  61.8   9.68
## 6     6  72.2   7.55
## 7     7  80.1   7.12
```

```
## 8      8  74.5    5.19
## 9      9  67.4    8.47
## 10     10 60.1    8.85
## 11     11 45.0   10.4
## 12     12 38.4    9.98
```

The standard deviation is a quantification of **spread** and **variability**. We see that the period in November, December, and January has the most variation in weather, so you can expect very different temperatures on different days in those months.

### LC 3.6 (Objective 1)

(LC3.6) What code would be required to get the mean and standard deviation temperature for each day in 2013 for NYC?

**Solution:**

```
summary(weather)
```

```
##      origin          year      month      day
## Length:26115      Min.   :2013      Min.   : 1.000      Min.   : 1.00
## Class :character  1st Qu.:2013      1st Qu.: 4.000      1st Qu.: 8.00
## Mode  :character  Median :2013      Median : 7.000      Median :16.00
##                               Mean  :2013      Mean   : 6.504      Mean   :15.68
##                               3rd Qu.:2013      3rd Qu.: 9.000      3rd Qu.:23.00
##                               Max.   :2013      Max.   :12.000      Max.   :31.00
##
##      hour      temp      dewp      humid
## Min.   : 0.00      Min.   : 10.94      Min.   : -9.94      Min.   : 12.74
## 1st Qu.: 6.00      1st Qu.: 39.92      1st Qu.:26.06      1st Qu.: 47.05
## Median :11.00      Median : 55.40      Median :42.08      Median : 61.79
## Mean   :11.49      Mean   : 55.26      Mean   :41.44      Mean   : 62.53
## 3rd Qu.:17.00      3rd Qu.: 69.98      3rd Qu.:57.92      3rd Qu.: 78.79
## Max.   :23.00      Max.   :100.04      Max.   :78.08      Max.   :100.00
##                               NA's   :1      NA's   :1      NA's   :1
##      wind_dir      wind_speed      wind_gust      precip
## Min.   : 0.0      Min.   : 0.000      Min.   :16.11      Min.   :0.000000
## 1st Qu.:120.0      1st Qu.: 6.905      1st Qu.:20.71      1st Qu.:0.000000
## Median :220.0      Median : 10.357      Median :24.17      Median :0.000000
## Mean   :199.8      Mean   : 10.518      Mean   :25.49      Mean   :0.004469
## 3rd Qu.:290.0      3rd Qu.: 13.809      3rd Qu.:28.77      3rd Qu.:0.000000
## Max.   :360.0      Max.   :1048.361      Max.   :66.75      Max.   :1.210000
## NA's   :460      NA's   :4      NA's   :20778
##      pressure      visib      time_hour
## Min.   : 983.8      Min.   : 0.000      Min.   :2013-01-01 01:00:00
## 1st Qu.:1012.9      1st Qu.:10.000      1st Qu.:2013-04-01 21:30:00
## Median :1017.6      Median :10.000      Median :2013-07-01 14:00:00
## Mean   :1017.9      Mean   : 9.255      Mean   :2013-07-01 18:26:37
## 3rd Qu.:1023.0      3rd Qu.:10.000      3rd Qu.:2013-09-30 13:00:00
## Max.   :1042.1      Max.   :10.000      Max.   :2013-12-30 18:00:00
## NA's   :2729
```

There is only one year 2013 so we don't need to group by it, we could but it would not change anything.

```
summary_temp_by_day <- weather %>%
  group_by(month, day) %>%
  summarize(
    mean = mean(temp, na.rm = TRUE),
    std_dev = sd(temp, na.rm = TRUE)
  )
```

## 'summarise()' has grouped output by 'month'. You can override using the  
## '.groups' argument.

```
head(summary_temp_by_day)
```

```
## # A tibble: 6 x 4
## # Groups:   month [1]
##   month   day mean std_dev
##   <int> <int> <dbl>   <dbl>
## 1     1     1  37.0     4.00
## 2     1     2  28.7     3.45
## 3     1     3  30.0     2.58
## 4     1     4  34.9     2.45
## 5     1     5  37.2     4.01
## 6     1     6  40.1     4.40
```

Note: `group_by(day)` is not enough, because `day` is a value between 1-31. We need to `group_by(year, month, day)` or `group_by(month, day)`.

### LC 3.7 (Objective 1)

(LC3.7) Recreate `by_monthly_origin`, but instead of grouping via `group_by(origin, month)`, group variables in a different order `group_by(month, origin)`. What differs in the resulting dataset?

**Solution:**

```
by_origin_monthly <- flights %>%
  group_by(origin, month) %>%
  summarize(count = n())
```

## 'summarise()' has grouped output by 'origin'. You can override using the  
## '.groups' argument.

```
head(by_origin_monthly)
```

```
## # A tibble: 6 x 3
## # Groups:   origin [1]
##   origin month count
##   <chr>   <int> <int>
## 1 EWR      1  9893
## 2 EWR      2  9107
## 3 EWR      3 10420
## 4 EWR      4 10531
## 5 EWR      5 10592
## 6 EWR      6 10175
```

```
by_monthly_origin <- flights %>%
  group_by(month, origin) %>%
  summarize(count = n())
```

## 'summarise()' has grouped output by 'month'. You can override using the  
## '.groups' argument.

```
head(by_monthly_origin)
```

```
## # A tibble: 6 x 3
## # Groups:   month [2]
##   month origin count
##   <int> <chr> <int>
## 1     1   EWR   9893
## 2     1   JFK   9161
## 3     1   LGA   7950
## 4     2   EWR   9107
## 5     2   JFK   8421
## 6     2   LGA   7423
```

In `by_monthly_origin` the `month` column is now first and the rows are sorted by `month` instead of `origin`. If you compare the values of `count` in `by_origin_monthly` and `by_monthly_origin` using the `View()` function, you'll see that the values are actually the same, just presented in a different order.

### LC 3.8 (Objective 1)

(LC3.8) How could we identify how many flights left each of the three airports for each `carrier`?

**Solution:** We could summarize the count from each airport and carrier using the `n()` function, which *counts rows*.

```
count_flights_by_airport <- flights %>%
  group_by(origin, carrier) %>%
  summarize(count = n())
```

```
head(count_flights_by_airport, n=10)
```

```
## # A tibble: 10 x 3
## # Groups:   origin [1]
##   origin carrier count
##   <chr> <chr> <int>
## 1 EWR    9E     1268
## 2 EWR    AA     3487
## 3 EWR    AS       714
## 4 EWR    B6     6557
## 5 EWR    DL     4342
## 6 EWR    EV    43939
## 7 EWR    MQ     2276
## 8 EWR    OO         6
## 9 EWR    UA    46087
## 10 EWR   US    4405
```

Note: the `n()` function counts rows, whereas the `sum(VARIABLE_NAME)` function sums all values of a certain numerical variable `VARIABLE_NAME`.

### LC 3.9 (Objective 1)

(LC3.9) How does the `filter` operation differ from a `group_by` followed by a `summarize`?

**Solution:**

- `filter` picks out rows from the original dataset without modifying them, whereas
- `group_by %>% summarize` computes summaries of numerical variables, and hence reports new values.

### LC 3.10 (Objective 2)

(LC3.10) What do positive values of the `gain` variable in `flights` correspond to? What about negative values? And what about a zero value?

**Solution:**

- Say a flight departed 20 minutes late, i.e. `dep_delay = 20`
- Then arrived 10 minutes late, i.e. `arr_delay = 10`.
- Then `gain = dep_delay - arr_delay = 20 - 10 = 10` is positive, so it “made up/gained time in the air.”
- 0 means the departure and arrival delay times were the same, so no time was made up in the air. We see in most cases that the `gain` is near 0 minutes.

### LC 3.11 (Objective 2)

(LC3.11) Could we create the `dep_delay` and `arr_delay` columns by simply subtracting `dep_time` from `sched_dep_time` and similarly for arrivals? Try the code out and explain any differences between the result and what actually appears in `flights`.

**Solution:** No because you can't do direct arithmetic on times. The difference in time between 12:03 and 11:59 is 4 minutes, but `1203-1159 = 44`. Plus there are time zones, departure and arrival times are in the local timezone, which cause problems with simple subtraction.

```
LC3.11<- flights %>%
  mutate(time_gain=dep_time-arr_time,gain = dep_delay - arr_delay) %>%
  select(air_time,dep_time,arr_time,time_gain,dep_delay,arr_delay, gain)
```

```
head(LC3.11)
```

```
## # A tibble: 6 x 7
##   air_time dep_time arr_time time_gain dep_delay arr_delay gain
##   <dbl>    <int>    <int>    <int>    <dbl>    <dbl> <dbl>
## 1     227      517      830     -313         2        11    -9
## 2     227      533      850     -317         4        20   -16
## 3     160      542      923     -381         2        33   -31
## 4     183      544     1004     -460        -1       -18    17
## 5     116      554      812     -258        -6       -25    19
## 6     150      554      740     -186        -4        12   -16
```

### LC 3.12 (Objective 2)

(LC3.12) What can we say about the distribution of `gain`? Describe it in a few sentences using a boxplot and the `gain_summary` data frame values.


**Solution:** We must create the data frame from the notes. We copied the code from the book, we had to copy and combine two chunks of code.

```
gain_summary <- flights %>%
  mutate(gain = dep_delay - arr_delay) %>%
  summarize(
    min = min(gain, na.rm = TRUE),
    q1 = quantile(gain, 0.25, na.rm = TRUE),
    median = quantile(gain, 0.5, na.rm = TRUE),
    q3 = quantile(gain, 0.75, na.rm = TRUE),
    max = max(gain, na.rm = TRUE),
    mean = mean(gain, na.rm = TRUE),
    sd = sd(gain, na.rm = TRUE),
    missing = sum(is.na(gain))
  )
```

```
gain_summary
```

```
## # A tibble: 1 x 8
##   min    q1 median    q3   max  mean    sd missing
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <int>
## 1  -196    -3     7    17   109  5.66  18.0   9430
```

```
flights %>%
  mutate(gain = dep_delay - arr_delay) %>%
  ggplot(aes(x=1,y=gain)) +
  geom_boxplot() +
  scale_x_continuous(breaks = NULL) +
  theme(axis.title.x = element_blank()) +
  theme_classic()
```



Most of the time the gain is a little above zero (the median is 7, meaning gain is above 0 at least 50% of the time) and between -50 and 50 minutes. There are some extreme cases however!

## Documenting software

- File creation date: 2022-06-16
- R version 4.1.3 (2022-03-10)
- ggplot2 package version: 3.3.6
- dplyr package version: 1.0.9
- nycflights13 package version: 1.0.2