

# Problem Set 06: Confidence Intervals

Professor Bradley Warner

June, 2022

**Documentation: None**

## Introduction

In this problem set we will use a small **sample** of data from the General Social Survey. The survey is designed to monitor changes in both social characteristics and attitudes. You will work with a **sample** from one neighborhood. The full neighborhood of **ALL individuals** is the population. For this problem set we do **not** know the **true population parameters** for any of the variables, because we do not have data on every person in the neighborhood. This is different from problem set 5 as we are using a sample of data.

## Setup

First load the necessary packages

```
# Recall that loading the tidyverse "umbrella" package loads ggplot2, dplyr, and  
# readr all at once. Feel free to load these packages any way you choose.  
library(tidyverse)  
library(moderndiver)  
library(infer)
```

Next load the data set, from where it is stored on the web:

```
gss_sample <- read_csv("https://docs.google.com/spreadsheets/d/e/2PACX-1vSypSoDCMH2N76Vo2dZRPkw2q3t1mbv...")
```

## Confidence intervals from a bootstrap resample

### Step 1: Take 1000 bootstrap resamples

The following code tells R to take 1000 bootstrap resamples from the `gss_sample` data. You can set the seed to whatever value you like!

```
set.seed(42)  
  
boot_samp_1000 <- gss_sample %>%  
  rep_sample_n(size = 100, reps = 1000, replace = TRUE)
```

Note a few important details about the `rep_sample_n` function, and bootstrap sampling in general:

- `size = 100` tells R that each bootstrap resample we take has 100 cases... the size of the original sample
- `reps = 1000` tells R to take 1000 bootstrap resamples (each of size 100).
- The `replace = TRUE` argument tells R that in each bootstrap resample, we can include a row from `gss_sample` multiple times. So if for instance, respondent # 12 is the first random resample taken here, respondent 12 is still available to be resampled **again** at random. Thus, some people may appear **multiple times** in our bootstrap resample, and some people from the original data set may not appear at all.
- We save the results in a data frame `boot_samp_1000`.

## Exercise 1

How many rows does `boot_samp_1000` have? **Why?**

**Answer: 100,000.** Because we took 1000 samples of 100 cases each, which is  $100,000 = \text{One hundred thousand}$ .

### Step 2: Calculate the bootstrap statistic

Let's say we want to use the bootstrap resample that we just generated to calculate a confidence interval for the population mean  $\mu_{tv}$  of `tvhours`. To do so, we need to know the sample mean  $\bar{x}$  of `tvhours` **for each of the 1000 bootstrap resamples**. In this case, the sample mean  $\bar{x}$  of `tvhours` for **each bootstrap resample** is our **BOOTSTRAP STATISTIC**. We can calculate that with two lines of code, like so:

```
boot_distrib_tv <- boot_samp_1000 %>%
  group_by(replicate) %>%
  summarize(boot_stat = mean(tvhours))
```

Note that:

- The `group_by()` argument tells R to take the sample mean of `tvhours` **separately** for each different `replicate` in the bootstrap resample.
- We put the sample mean for each bootstrap resample in a column called `boot_stat`

This is the bootstrap distribution for the mean of `tvhours`!


Take a look at the `boot_distrib_tv` we just created in RStudio's data viewer.

## Exercise 2

How many values of the bootstrap statistic `boot_stat` are there? Please explain **why** there are this many values of the bootstrap statistic.

**Answer: there are 1000 `boot_stat` values, because we took 1000 bootstrap resamples, and calculated the sample mean for each sample. This equates to 1000 bootstrap statistics.**

**Visualizing the bootstrap distribution** The bootstrap distribution is shown in the figure below. This is a histogram of the `boot_stat` values from `boot_distrib_tv`.



### Step 3: CI from a bootstrap resample

**a) CI using the standard error** We can now use the bootstrap distribution for the sample mean `tvhours`  $\bar{x}$  to calculate a 95% confidence interval for the population mean `tvhours`  $\mu_{tv}$ , using the standard error method. This method assumes the bootstrap distribution is approximately normal. This interval is calculated as

$$\text{mean} \pm 1.96 \cdot SD$$

The 1.96 comes from the standard normal distribution.

```
qnorm(.975)
```

```
## [1] 1.959964
```

- the mean here would be the mean of the bootstrap distribution
- the SD here is the standard deviation of the bootstrap distribution, which recall has a special name: the **standard error**.

We can thus apply the standard error method, like so:

```
boot_distrib_tv %>%
  summarize(mean = mean(boot_stat),
            se = sd(boot_stat),
            lower_ci = mean - (qnorm(.975) * se),
            upper_ci = mean + (qnorm(.975) * se))
```

```
## # A tibble: 1 x 4
##   mean    se lower_ci upper_ci
##   <dbl> <dbl>   <dbl>   <dbl>
## 1  3.14 0.365    2.43    3.86
```

**b) CI using percentile method** You can also calculate a 95% confidence interval using the percentile method. The logic goes like this:

Since our bootstrap resample had 1000 values of `boot_stat`:

- 950 of the `boot_stat` values fall **inside** this 95% confidence interval, i.e. 95%
- 25 values fall **below** it. i.e. the lower 2.5%
- 25 values fall **above** it. i.e. the higher 2.5%

totaling 100%. We can use **the quantiles** of the bootstrap distribution to find these values like so:

```
prct_ci_tv <- boot_distrib_tv %>%  
  summarize(lower_ci = quantile(boot_stat, 0.025),  
            upper_ci = quantile(boot_stat, 0.975))  
  
prct_ci_tv
```

```
## # A tibble: 1 x 2  
##   lower_ci upper_ci  
##   <dbl>    <dbl>  
## 1     2.51     3.89
```


This method

- Asks R to identify the 0.025 quantile of the bootstrap sample means... this is the value **below** which **2.5% of the values of `boot_stat`** fall (or 25 cases in this example...  $25/1000 = 0.025$ )
- Asks R to identify the 0.975 quantile for the bootstrap sample means... this is the value **above** which the other **2.5% of the values of `boot_stat`** fall (or 25 cases in this example  $975/1000 = 0.975$ )
- The middle 95% of the values fall between these two quantiles

Based on these results, we are 95% confident that the **true mean hours of TV watched  $\mu_{tv}$  in the population** is between the upper and lower CI we just calculated.

## Visualizing the Confidence interval

The bootstrap distribution and the 95% confidence intervals we just calculated are shown in the figure below. This is a histogram of the `boot_stat` values from `boot_distrib_tv`. The green line is the lower bound of the 95% CI, and the blue line is the upper bound. 950 of the 1000 bootstrap resamples had a mean for `tvhours` that fell **between** the green and blue lines... 25 of the samples had a mean above the blue line, and 25 of the samples had a mean below the green line.



### Exercise 3

If we calculated a **90% confidence interval** for the mean of `tvhours` using this same bootstrap resample and the percentile method, roughly how many of the 1000 values of `tv_mean` would fall between the green and blue lines?

Answer: 900

### Exercise 4

Use the bootstrap resampling distribution for `tvhours` generated above (`boot_distrib_tv`) and the **percentile** method to calculate a 99% confidence interval for the mean `tvhours`.

```
boot_distrib_tv %>%
  summarize(lower_ci = quantile(boot_stat, 0.005),
            upper_ci = quantile(boot_stat, 0.995))
```

```
## # A tibble: 1 x 2
##   lower_ci upper_ci
##   <dbl>    <dbl>
## 1     2.32     4.13
```

### Exercise 5

Which confidence interval is **WIDER**: the 95% confidence interval or the 99% confidence interval for the population mean `tvhours`  $\mu_{tv}$ ? Why is this the case? Answer in terms of the trade-off between confidence level and confidence interval width.

Answer: Conceptually, the 99% CI is wider, because we are casting a wider net. In other words, to have a higher degree of certainty we captured the population mean  $\mu_{tv}$ , we need a wider range of values contained in our confidence interval. Also, more of the `boot_stat` values are included in a 99% CI (990) as compared to a 95% CI (950)

## Finding Confidence Intervals Using Infer

As we learning in the reading, the `infer` package provides a way to find confidence intervals. We used this package in many of our in-class notes.

### Exercise 6

Use the `infer` package and the the functions `specify()`, `generate()`, and `calculate()` to generate a bootstrap distribution for the variable `race`. Set the seed to 504. Success is POC. Use 1000 bootstrap replications.

**Answer:**

```
set.seed(504)
infer_boot<-gss_sample %>%
  specify(formula=race~NULL,succes="POC") %>%
  generate(reps = 1000, type = "bootstrap") %>%
  calculate(stat="prop")
```

```
head(infer_boot)
```

```
## Response: race (factor)
## # A tibble: 6 x 2
##   replicate stat
##   <int> <dbl>
## 1      1  0.19
## 2      2  0.33
## 3      3  0.26
## 4      4  0.21
## 5      5  0.21
## 6      6  0.16
```

### Exercise 7

Calculate a 90% confidence interval for the proportion of respondents that identify as POC using the **standard error** method.

**Answer:**

```
gss_sample %>%
  summarize(prob=mean(race=="POC"))
```

```
## # A tibble: 1 x 1
##   prob
##   <dbl>
## 1  0.24
```

```
get_ci(infer_boot,level=.90,type="se",point_estimate = 0.24)
```

```
## # A tibble: 1 x 2
##   lower_ci upper_ci
##   <dbl>    <dbl>
## 1    0.168    0.312
```

## Exercise 8

Calculate a 90% confidence interval for the proportion of respondents that identify as POC using the **percentile** method.

**Answer:**

```
get_ci(infer_boot, level=.90, type="percentile")
```

```
## # A tibble: 1 x 2
##   lower_ci upper_ci
##   <dbl>    <dbl>
## 1    0.17    0.31
```

## Exercise 9

Interpret the 90% **percentile** confidence interval.

**Answer:**


We are 90% confident that the true proportion of respondents that identify as a person of color is between 0.168311 and 0.311689.

## Exercise 10

Create a visualization, using the **infer** package, of the bootstrap sampling distribution. Discuss if it is appropriate to use the standard error method.

**Answer:**

```
visualize(infer_boot)
```



The distribution appears to be approximately normal so the use of the standard error method seems appropriate.

---

## Documenting software

- File creation date: 2022-06-13
- R version 4.1.3 (2022-03-10)
- **tidyverse** package version: 1.3.1
- **moderndive** package version: 0.5.4
- **infer** package version: 1.0.0